

SGD_Assignment

March 22, 2020

```
[1]: import numpy as np
import pandas as pd
from tqdm.notebook import tqdm
from sklearn.datasets import make_classification

[2]: X, y = make_classification(n_samples=50000, n_features=15, n_informative=10,
    ↪n_redundant=5,
                                n_classes=2, weights=[0.7], class_sep=0.7,
    ↪random_state=15)

[3]: X.shape, y.shape

[3]: ((50000, 15), (50000,))

[4]: from sklearn.model_selection import train_test_split

[5]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25,
    ↪random_state=15)

[6]: X_train.shape, y_train.shape, X_test.shape, y_test.shape

[6]: ((37500, 15), (37500,)), (12500, 15), (12500,))

[7]: from sklearn import linear_model

[8]: # alpha : float
# Constant that multiplies the regularization term.

# eta0 : double
# The initial learning rate for the 'constant', 'invscaling' or 'adaptive'
    ↪schedules.

clf = linear_model.SGDClassifier(eta0=0.0001, alpha=0.0001, loss='log',
    ↪random_state=15, penalty='l2', tol=1e-3, verbose=2, learning_rate='constant')
clf
```

```
[8]: SGDClassifier(alpha=0.0001, average=False, class_weight=None,
                  early_stopping=False, epsilon=0.1, eta0=0.0001,
                  fit_intercept=True, l1_ratio=0.15, learning_rate='constant',
                  loss='log', max_iter=1000, n_iter_no_change=5, n_jobs=None,
                  penalty='l2', power_t=0.5, random_state=15, shuffle=True,
                  tol=0.001, validation_fraction=0.1, verbose=2, warm_start=False)
```

```
[9]: clf.fit(X=X_train, y=y_train)
```

```
-- Epoch 1
Norm: 0.77, NNZs: 15, Bias: -0.316653, T: 37500, Avg. loss: 0.455552
Total training time: 0.02 seconds.
-- Epoch 2
Norm: 0.91, NNZs: 15, Bias: -0.472747, T: 75000, Avg. loss: 0.394686
Total training time: 0.04 seconds.
-- Epoch 3
Norm: 0.98, NNZs: 15, Bias: -0.580082, T: 112500, Avg. loss: 0.385711
Total training time: 0.06 seconds.
-- Epoch 4
Norm: 1.02, NNZs: 15, Bias: -0.658292, T: 150000, Avg. loss: 0.382083
Total training time: 0.08 seconds.
-- Epoch 5
Norm: 1.04, NNZs: 15, Bias: -0.719528, T: 187500, Avg. loss: 0.380486
Total training time: 0.10 seconds.
-- Epoch 6
Norm: 1.05, NNZs: 15, Bias: -0.763409, T: 225000, Avg. loss: 0.379578
Total training time: 0.12 seconds.
-- Epoch 7
Norm: 1.06, NNZs: 15, Bias: -0.795106, T: 262500, Avg. loss: 0.379150
Total training time: 0.14 seconds.
-- Epoch 8
Norm: 1.06, NNZs: 15, Bias: -0.819925, T: 300000, Avg. loss: 0.378856
Total training time: 0.17 seconds.
-- Epoch 9
Norm: 1.07, NNZs: 15, Bias: -0.837805, T: 337500, Avg. loss: 0.378585
Total training time: 0.19 seconds.
-- Epoch 10
Norm: 1.08, NNZs: 15, Bias: -0.853138, T: 375000, Avg. loss: 0.378630
Total training time: 0.21 seconds.
Convergence after 10 epochs took 0.22 seconds
```

```
[9]: SGDClassifier(alpha=0.0001, average=False, class_weight=None,
                  early_stopping=False, epsilon=0.1, eta0=0.0001,
                  fit_intercept=True, l1_ratio=0.15, learning_rate='constant',
                  loss='log', max_iter=1000, n_iter_no_change=5, n_jobs=None,
                  penalty='l2', power_t=0.5, random_state=15, shuffle=True,
                  tol=0.001, validation_fraction=0.1, verbose=2, warm_start=False)
```

```
[10]: clf.coef_, clf.coef_.shape, clf.intercept_
```

```
[10]: (array([[ -0.42336692,  0.18547565, -0.14859036,  0.34144407, -0.2081867 ,
              0.56016579, -0.45242483, -0.09408813,  0.2092732 ,  0.18084126,
              0.19705191,  0.00421916, -0.0796037 ,  0.33852802,  0.02266721]]),
      (1, 15),
      array([ -0.8531383]))
```

0.1 Implement Logistic Regression with L2 regularization Using SGD: without using sklearn

0.1.1 Instructions

- Load the datasets(train and test) into the respective arrays
- Initialize the weight_vector and intercept term randomly
- Calculate the initial log loss for the train and test data with the current weight and intercept and store it in a list
- for each epoch:
 - for each batch of data points in train: (keep batch size=1)
 - * calculate the gradient of loss function w.r.t each weight in weight vector
 - * Calculate the gradient of the intercept check this
 - * Update weights and intercept (check the equation number 32 in the above mentioned pdf): $w^{(t+1)} \leftarrow (1 - \frac{1}{N})w^{(t)} + x_n(y_n - ((w^{(t)})^T x_n + b^t))$ $b^{(t+1)} \leftarrow (b^t + (y_n - ((w^{(t)})^T x_n + b^t)))$
 - * calculate the log loss for train and test with the updated weights (you can check the python assignment 10th question)
 - * And if you wish, you can compare the previous loss and the current loss, if it is not updating, then you can stop the training
 - * append this loss in the list (this will be used to see how loss is changing for each epoch after the training is over)
- Plot the train and test loss i.e on x-axis the epoch number, and on y-axis the loss
- GOAL: compare your implementation and SGDClassifier's the weights and intercept, make sure they are as close as possible i.e difference should be in terms of 10^{-3}

```
[11]: w = np.zeros_like(X_train[0])
      b = 0
      eta0 = 0.0001
      alpha = 0.0001
      N = len(X_train)
```

```
[12]: #We already have data loaded to X_train, X_test, y_train, y_test
      #Vales of Weight vector and intercept term are initially set to 0
```

```
[13]: #Sigmoid function
      def sigmoid(w,x,b):
          val=1/(1+np.exp(-(np.dot(x,w.T)+b)))
```

```
return val
```

```
[14]: #Log Loss function
# Log loss= -((1/N)*[(y*log(p))+((1-y)*log(1-p))])
def log_loss(w,x,b,y):
    sigmoid_val=sigmoid(w,x,b)
    first_term=y*np.log(sigmoid_val)
    second_term=(1-y)*np.log(1-sigmoid_val)
    val=-((first_term+second_term)/(len(x)))
    return np.sum(val)
```

```
[15]: # write your code to implement SGD as per the above instructions
# please choose the number of iterations on your own
```

```
[31]: def SGD(w,b,epoch,x_train,y_train,x_test,y_test):
    log_loss_train=[]
    log_loss_test=[]
    N=len(x_train)
    #Calculate initial log loss and storing the values
    log_loss_train.append(log_loss(w,x_train,b,y_train))
    log_loss_test.append(log_loss(w,x_test,b,y_test))

    for epo in tqdm(range(epoch)):
        for i in range(N):
            w= ((1-(alpha*eta0)/
↪N)*w)+((alpha*x_train[i])*(y_train[i]-sigmoid(w,x_train[i],b)))
            b= (b+(alpha*(y_train[i]-sigmoid(w,x_train[i],b))))
            log_val_train=log_loss(w,x_train,b,y_train)
            log_val_test=log_loss(w,x_test,b,y_test)
            log_loss_train.append(log_val_train)
            log_loss_test.append(log_val_test)
            if abs((np.sum(log_loss_train[epo+1]-log_loss_train[epo])))<=0.0001:
                print('Updation is completed in {} epochs'.format(epo+1))
                break

    return w,b,log_loss_train,log_loss_test,epo+1
```

```
[32]: epochs=20
updated_w,updated_b,train_log_loss,test_log_loss,convergence_epoch=SGD(w,b,epochs,X_train,y_train,
↪X_test, y_test)
```

```
HBox(children=(IntProgress(value=0, max=20), HTML(value='')))
```

Updation is completed in 10 epochs

```
[34]: print(updated_w)
      print(updated_b)
```

```
[-0.42315311  0.19095979 -0.14588118  0.33814991 -0.21196623  0.56525978
 -0.44538357 -0.09171679  0.21795314  0.16977398  0.19522044  0.00229554
 -0.07781461  0.33882618  0.02214234]
-0.8500967712837224
```

```
[35]: # these are the results we got after we implemented sgd and found the optimal_
      ↪ weights and intercept
      updated_w-clf.coef_, updated_b-clf.intercept_
```

```
[35]: (array([[ 0.0002138 ,  0.00548413,  0.00270918, -0.00329416, -0.00377953,
                0.00509399,  0.00704126,  0.00237134,  0.00867994, -0.01106728,
               -0.00183147, -0.00192361,  0.00178909,  0.00029817, -0.00052487]]),
      array([0.00304153]))
```

```
[36]: def pred(w,b, X):
      N = len(X)
      predict = []
      for i in range(N):
          if sigmoid(w, X[i], b) >= 0.5: # sigmoid(w,x,b) returns 1/
          ↪ (1+exp(-(dot(x,w)+b)))
              predict.append(1)
          else:
              predict.append(0)
      return np.array(predict)
      print(1-np.sum(y_train - pred(updated_w,updated_b,X_train))/len(X_train))
      print(1-np.sum(y_test - pred(updated_w,updated_b,X_test))/len(X_test))
```

```
0.95536
0.95296
```

```
[43]: import matplotlib.pyplot as plt
      #Plotting without initial log loss as it is very high and makes the two curve_
      ↪ in the plot to be not clear
      plt.plot(range(1,convergence_epoch+1),train_log_loss[1:], label='Training Data_
      ↪ log loss', color='blue')
      plt.plot(range(1,convergence_epoch+1),test_log_loss[1:], label='Test Data log_
      ↪ loss', color='green')

      plt.ylabel('Log Loss')
      plt.xlabel('epochs')
      plt.title('Log loss')
      plt.legend()
      plt.show()
```

