

Transfer Learning Assignment

Download all the data in this [rar_file](#) , it contains all the data required for the assignment. When you unrar the file you'll get the files in the following format:

path/to/the/image.tif,category

where the categories are numbered 0 to 15, in the following order:

- 0 letter
- 1 form
- 2 email
- 3 handwritten
- 4 advertisement
- 5 scientific report
- 6 scientific publication
- 7 specification
- 8 file folder
- 9 news article
- 10 budget
- 11 invoice
- 12 presentation
- 13 questionnaire
- 14 resume
- 15 memo

There is a file named as 'labels_final.csv' , it consists of two columns. First column is path which is the required path to the images and second is the class label.

```
In [ ]: #the dataset that you are dealing with is quite large 3.7 GB and hence there are two methods
# Method 1- you can use gdown module to get the data directly from Google drive to
# the syntax is as follows !gdown --id file_id , for ex - running the below cell will
```

```
In [ ]: !gdown --id 1Z4TyI7FcFVEx8qdl4j09qxvxaqLSqoEu
```

```
In [ ]: # Method -2 you can also import the data using wget function
#https://www.youtube.com/watch?v=BPUfVq7RaY8
```

```
In [ ]: #unrar the file
#get_ipython().system_raw("unrar x rvL-cdip.rar")
```

2. On this image data, you have to train 3 types of models as given below You have to split the data into Train and Validation data.

```
In [1]: #import all the required libraries
```

```

import tensorflow as tf
import os
import numpy as np
import pandas as pd
from tensorflow.keras.layers import MaxPool2D, Conv2D, Flatten, Dense
from tensorflow.keras import Model
from tensorflow.keras.callbacks import ModelCheckpoint, EarlyStopping
import datetime
import os
import random as rn
from tensorflow.keras.utils import plot_model

%load_ext tensorboard

```

```
In [2]: df=pd.read_csv('labels_final.csv',dtype=str)
```

```
In [3]: df.head()
```

```
Out[3]:
```

	path	label
0	imagesv/v/o/h/voh71d00/509132755+-2755.tif	3
1	imagesl/l/x/t/lxt19d00/502213303.tif	3
2	imagesx/x/e/d/xed05a00/2075325674.tif	2
3	imageso/o/j/b/ojb60d00/517511301+-1301.tif	3
4	imagesq/q/z/k/qzk17e00/2031320195.tif	7

```
In [4]: df.shape
```

```
Out[4]: (48000, 2)
```

```
In [5]: # https://vijayabhaskar96.medium.com/tutorial-on-keras-flow-from-dataframe-1fd4493d
ImageFlow = tf.keras.preprocessing.image.ImageDataGenerator(rescale=1./255.,validat
```

```

train_generator=ImageFlow.flow_from_dataframe(dataframe=df,
directory="./data_final/",
x_col="path",
y_col="label",
subset="training",
batch_size=64,
seed=0,
shuffle=True,
class_mode="categorical",
target_size=(128,128))

validation_generator=ImageFlow.flow_from_dataframe(dataframe=df,
directory="./data_final/",
x_col="path",
y_col="label",
subset="validation",
batch_size=64,

```

```
seed=0,  
shuffle=True,  
class_mode="categorical",  
target_size=(128,128))
```

Found 36000 validated image filenames belonging to 16 classes.
Found 12000 validated image filenames belonging to 16 classes.

```
In [6]: train_steps_per_epoch = np.ceil(36000/64)  
validation_steps_per_epoch = np.ceil(12000/64)
```

3. Try not to load all the images into memory, use the generators that we have given the reference notebooks to load the batch of images only during the train data.

or you can use this method also <https://medium.com/@vijayabhaskar96/tutorial-on-keras-imagedatagenerator-with-flow-from-dataframe-8bd5776e45c1>

<https://medium.com/@vijayabhaskar96/tutorial-on-keras-flow-from-dataframe-1fd4493d237c>

Note- In the reference notebook you were dealing with jpg images, in the given dataset you are dealing with tiff images. Imagedatagenrator works with both type of images. If you want to use custom data pipeline then you have to convert your tiff images to jpg images.

4. You are free to choose Learning rate, optimizer, loss function, image augmentation, any hyperparameters. but you have to use the same architecture what we are asking below.
5. Use tensorboard for every model and analyse your gradients. (you need to upload the screenshots for each model for evaluation)
6. You can check about Transfer Learning in this link - <https://blog.keras.io/building-powerful-image-classification-models-using-very-little-data.html>

<https://www.appliedaicourse.com/lecture/11/applied-machine-learning-online-course/3426/code-example-cats-vs-dogs/8/module-8-neural-networks-computer-vision-and-deep-learning>

7. Do print model.summary() and draw model_plots for each of the model.

Model-1

1. Use **VGG-16** pretrained network without Fully Connected layers and initialize all the weights with Imagenet trained weights.
2. After VGG-16 network without FC layers, add a new Conv block (1 Conv layer and 1 Maxpooling), 2 FC layers and an output layer to classify 16 classes. You are free to choose any hyperparameters/parameters of conv block, FC layers, output layer.
3. Final architecture will be **INPUT --> VGG-16 without Top**

layers(FC) --> Conv Layer --> Maxpool Layer --> 2 FC layers --> Output Layer

4. Print model.summary() and plot the architecture of the model.

[Reference for plotting model](#)

5. Train only new Conv block, FC layers, output layer. Don't train the VGG-16 network.

```
In [7]: os.environ['PYTHONHASHSEED'] = '0'
```

```
##https://keras.io/getting-started/faq/#how-can-i-obtain-reproducible-results-using  
## Have to clear the session. If you are not clearing, Graph will create again and  
## Variables will also set to some value from before session  
tf.keras.backend.clear_session()
```

```
## Set the random seed values to regenerate the model.  
np.random.seed(0)  
rn.seed(0)  
tf.random.set_seed(0)
```

```
In [9]: vgg_model = tf.keras.applications.VGG16(  
        include_top=False,  
        weights="imagenet",  
        input_shape=(128,128,3),  
    )  
vgg_model.summary()
```

Model: "vgg16"

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	[(None, 128, 128, 3)]	0
=====		
block1_conv1 (Conv2D)	(None, 128, 128, 64)	1792
block1_conv2 (Conv2D)	(None, 128, 128, 64)	36928
=====		
block1_pool (MaxPooling2D)	(None, 64, 64, 64)	0
=====		
block2_conv1 (Conv2D)	(None, 64, 64, 128)	73856
block2_conv2 (Conv2D)	(None, 64, 64, 128)	147584
=====		
block2_pool (MaxPooling2D)	(None, 32, 32, 128)	0
=====		
block3_conv1 (Conv2D)	(None, 32, 32, 256)	295168
block3_conv2 (Conv2D)	(None, 32, 32, 256)	590080
block3_conv3 (Conv2D)	(None, 32, 32, 256)	590080
=====		
block3_pool (MaxPooling2D)	(None, 16, 16, 256)	0
=====		
block4_conv1 (Conv2D)	(None, 16, 16, 512)	1180160
block4_conv2 (Conv2D)	(None, 16, 16, 512)	2359808
block4_conv3 (Conv2D)	(None, 16, 16, 512)	2359808
=====		
block4_pool (MaxPooling2D)	(None, 8, 8, 512)	0
=====		
block5_conv1 (Conv2D)	(None, 8, 8, 512)	2359808
block5_conv2 (Conv2D)	(None, 8, 8, 512)	2359808
block5_conv3 (Conv2D)	(None, 8, 8, 512)	2359808
=====		
block5_pool (MaxPooling2D)	(None, 4, 4, 512)	0
=====		
Total params: 14,714,688		
Trainable params: 14,714,688		
Non-trainable params: 0		

```
In [9]: #https://blog.keras.io/building-powerful-image-classification-models-using-very-lit
for layer in vgg_model.layers:
    layer.trainable = False
input_from_vgg = vgg_model.output
conv1 = Conv2D(filters = 512, kernel_size = (3,3), padding = 'same',
               activation = 'relu', kernel_initializer='HeUniform') (input_from_vgg)
max_pool1 = MaxPool2D(pool_size=(2, 2)) (conv1)
flat = Flatten() (max_pool1)
dense_1 = Dense(256, activation = 'relu', kernel_initializer='HeUniform') (flat)
```

```
dense_2 = Dense(128, activation = 'relu', kernel_initializer='HeUniform')(dense_1)
output = Dense(16, activation='softmax')(dense_2)

model_one = Model(inputs = vgg_model.input, outputs = output)
model_one.summary()
```

Model: "model"

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	[(None, 128, 128, 3)]	0
block1_conv1 (Conv2D)	(None, 128, 128, 64)	1792
block1_conv2 (Conv2D)	(None, 128, 128, 64)	36928
block1_pool (MaxPooling2D)	(None, 64, 64, 64)	0
block2_conv1 (Conv2D)	(None, 64, 64, 128)	73856
block2_conv2 (Conv2D)	(None, 64, 64, 128)	147584
block2_pool (MaxPooling2D)	(None, 32, 32, 128)	0
block3_conv1 (Conv2D)	(None, 32, 32, 256)	295168
block3_conv2 (Conv2D)	(None, 32, 32, 256)	590080
block3_conv3 (Conv2D)	(None, 32, 32, 256)	590080
block3_pool (MaxPooling2D)	(None, 16, 16, 256)	0
block4_conv1 (Conv2D)	(None, 16, 16, 512)	1180160
block4_conv2 (Conv2D)	(None, 16, 16, 512)	2359808
block4_conv3 (Conv2D)	(None, 16, 16, 512)	2359808
block4_pool (MaxPooling2D)	(None, 8, 8, 512)	0
block5_conv1 (Conv2D)	(None, 8, 8, 512)	2359808
block5_conv2 (Conv2D)	(None, 8, 8, 512)	2359808
block5_conv3 (Conv2D)	(None, 8, 8, 512)	2359808
block5_pool (MaxPooling2D)	(None, 4, 4, 512)	0
conv2d (Conv2D)	(None, 4, 4, 512)	2359808
max_pooling2d (MaxPooling2D)	(None, 2, 2, 512)	0
flatten (Flatten)	(None, 2048)	0
dense (Dense)	(None, 256)	524544
dense_1 (Dense)	(None, 128)	32896
dense_2 (Dense)	(None, 16)	2064
=====		
Total params: 17,634,000		
Trainable params: 2,919,312		

Non-trainable params: 14,714,688

```
In [10]: filepath="model_save/best_model_one.hdf5"
checkpoint = ModelCheckpoint(filepath=filepath, monitor='val_accuracy', verbose=1,
earlystop = EarlyStopping(monitor='val_accuracy', min_delta=0.01, patience=2, verbo
tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir= ("logs/fits/"+dateti
histogram_freq=1,write_graph=
call_back_list = [ earlystop, checkpoint, tensorboard_callback]
```

```
In [11]: model_one.compile(optimizer='adam', loss='categorical_crossentropy',metrics=['accu
model_one.fit( train_generator,steps_per_epoch = train_steps_per_epoch,epochs=3,
validation_data = validation_generator, validation_steps =
callbacks=[call_back_list])
```

Epoch 1/3

563/563 [=====] - 178s 300ms/step - loss: 1.5444 - accuracy: 0.5238 - val_loss: 1.2278 - val_accuracy: 0.6213

Epoch 00001: val_accuracy improved from -inf to 0.62133, saving model to model_save\best_model_one.hdf5

Epoch 2/3

563/563 [=====] - 150s 266ms/step - loss: 1.1656 - accuracy: 0.6388 - val_loss: 1.1875 - val_accuracy: 0.6358

Epoch 00002: val_accuracy improved from 0.62133 to 0.63583, saving model to model_save\best_model_one.hdf5

Epoch 3/3

563/563 [=====] - 149s 265ms/step - loss: 1.0246 - accuracy: 0.6807 - val_loss: 1.1538 - val_accuracy: 0.6527

Epoch 00003: val_accuracy improved from 0.63583 to 0.65267, saving model to model_save\best_model_one.hdf5


Out[11]: <keras.callbacks.History at 0x1fe081a2c40>

```
In [52]: %tensorboard --logdir logs/fits
```


Reusing TensorBoard on port 6006 (pid 880), started 0:01:06 ago. (Use '!kill 880' to kill it.)

model_one Tensorboard accuracy and loss

epoch_accuracy
tag: epoch_accuracy

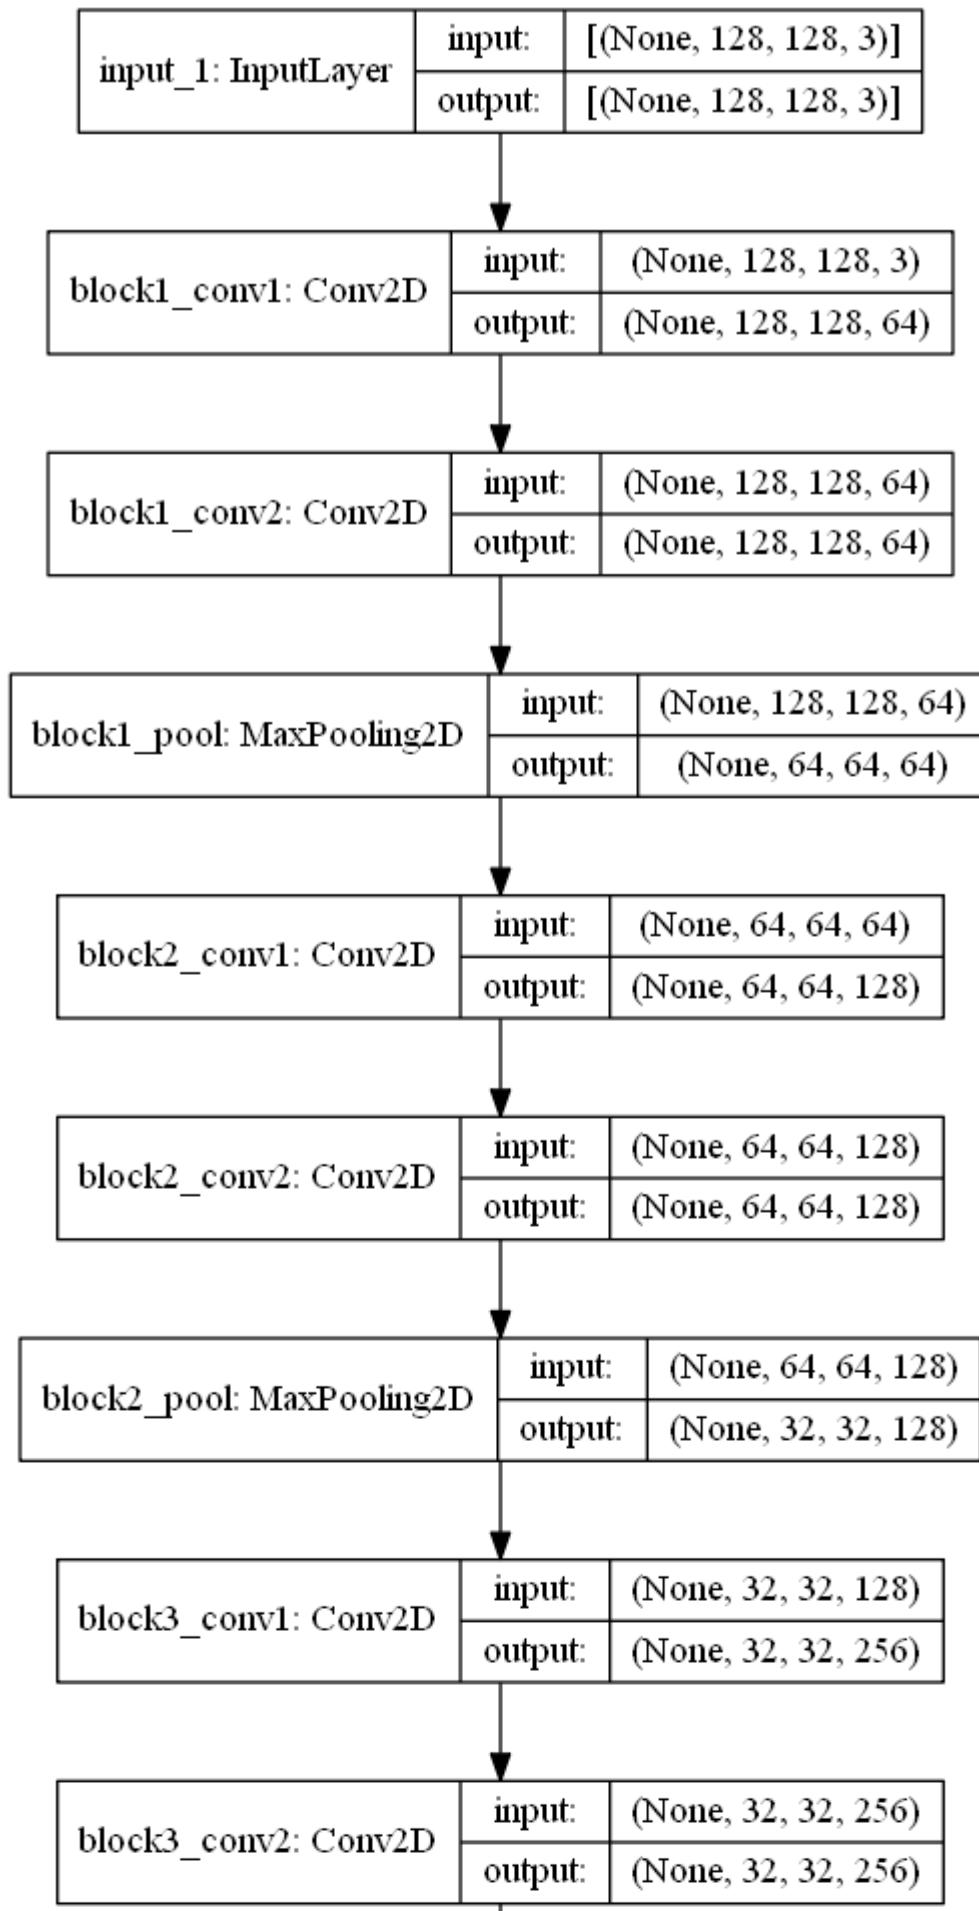


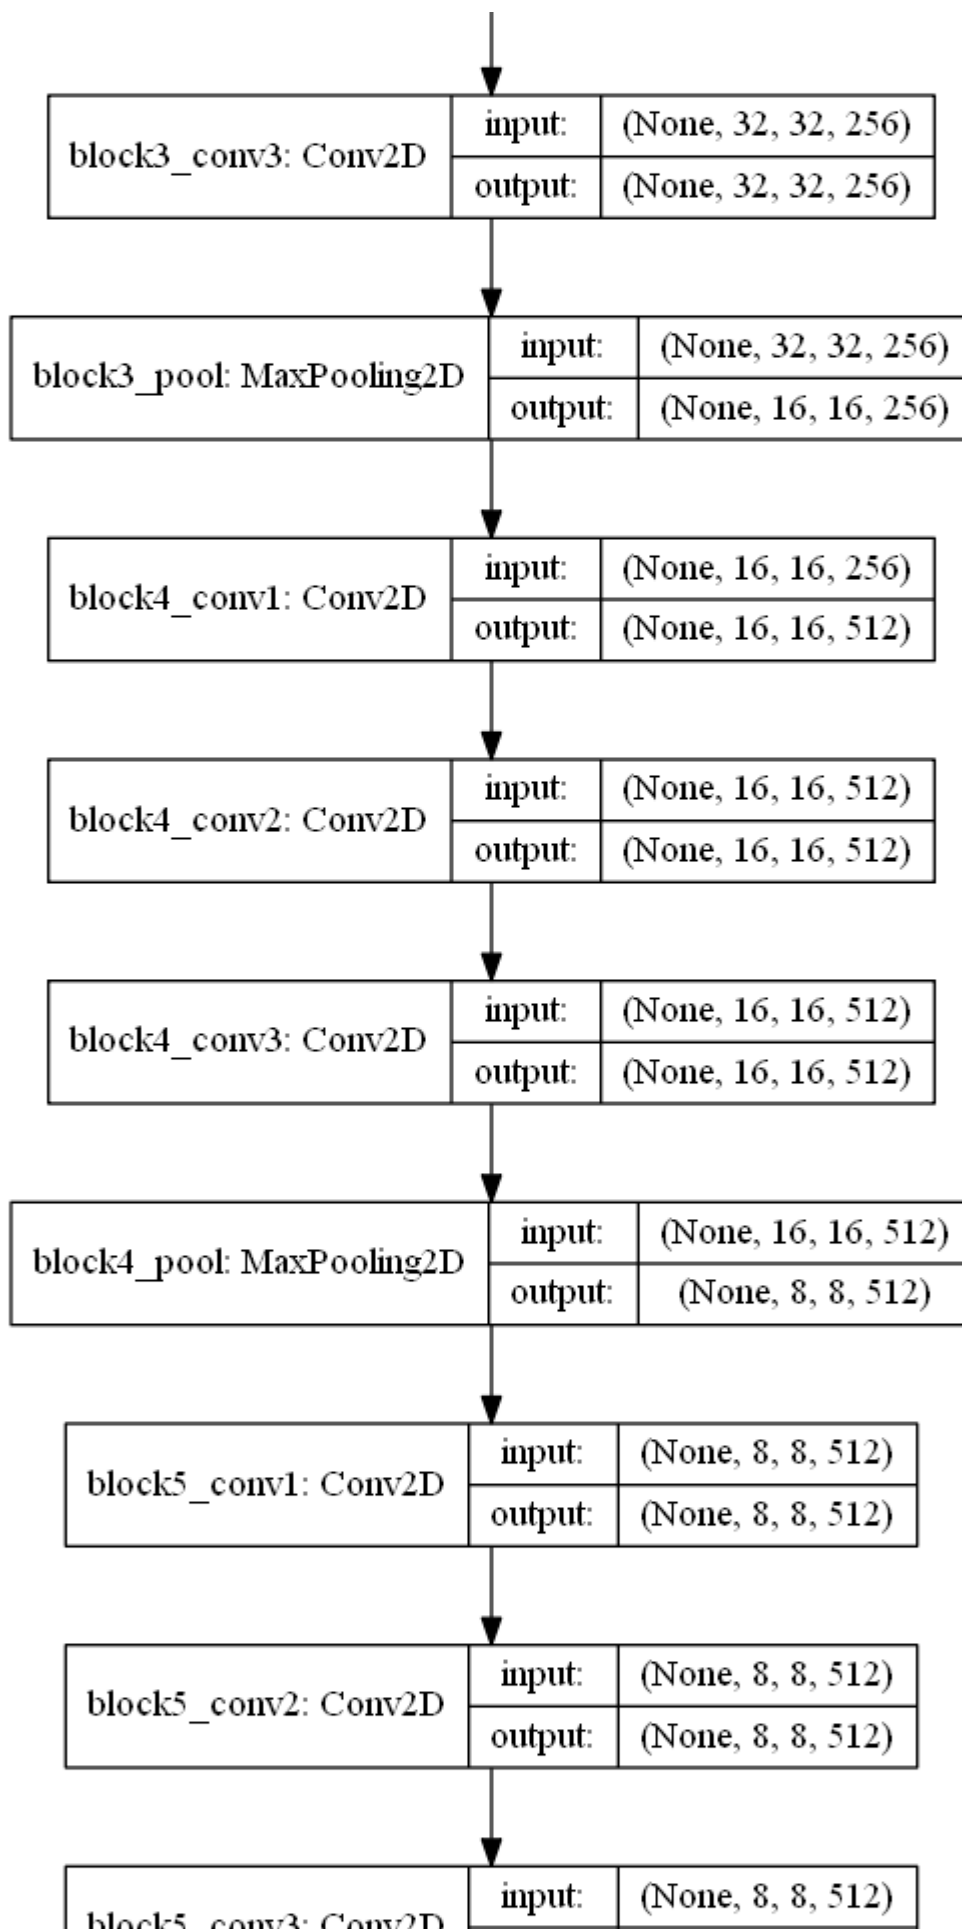
epoch_loss
tag: epoch_loss

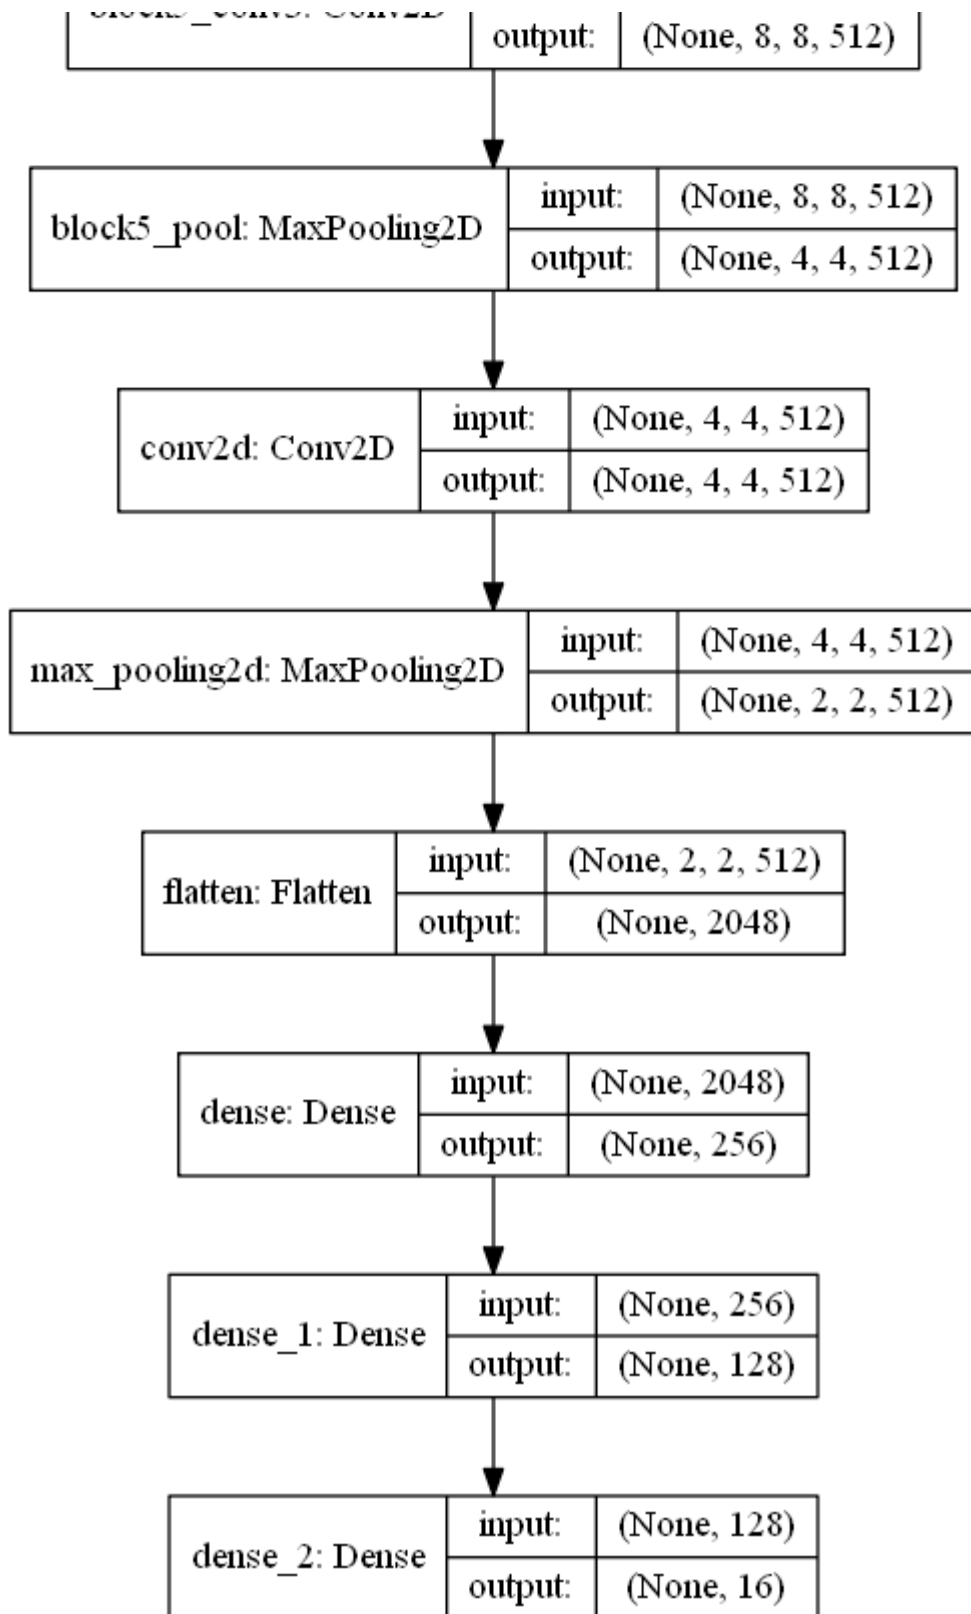


```
In [16]: plot_model(model_one, to_file='model_one.png', show_shapes=True)
```

Out[16]:







Model-2

1. Use [VGG-16](#) pretrained network without Fully Connected layers and initialize all the weights with Imagenet trained weights.
2. After VGG-16 network without FC layers, don't use FC layers, use

conv layers only as Fully connected layer. Any FC layer can be converted to a CONV layer. This conversion will reduce the No of Trainable parameters in FC layers.

For example, an FC layer with $K=4096$ that is looking at some input volume of size $7 \times 7 \times 512$ can be equivalently expressed as a CONV layer with $F=7, P=0, S=1, K=4096$.

In other words, we are setting the filter size to be exactly the size of the input volume, and hence the output will simply be $1 \times 1 \times 4096$ since only a single depth column "fits" across the input volume, giving identical result as the initial FC layer. You can refer [this](#) link to better understanding of using Conv layer in place of fully connected layers.

3. Final architecture will be VGG-16 without FC layers(without top), 2 Conv layers identical to FC layers, 1 output layer for 16 class classification. **INPUT --> VGG-16 without Top layers(FC) --> 2 Conv Layers identical to FC --> Output Layer**

4. 4. Print model.summary() and plot the architecture of the model. [Reference for plotting model](#)

5. Train only last 2 Conv layers identical to FC layers, 1 output layer. Don't train the VGG-16 network.

```
In [12]: tf.keras.backend.clear_session()
```

```
In [13]: #https://blog.keras.io/building-powerful-image-classification-models-using-very-lit
for layer in vgg_model.layers:
    layer.trainable = False
input_from_vgg = vgg_model.output

conv_1 = Conv2D(filters = 2048, kernel_size = (4,4), padding = 'valid',
                activation = 'relu', kernel_initializer='HeUniform')(input_from_vgg

conv_2 = Conv2D(filters = 2048, kernel_size = (1,1), padding = 'valid',
                activation = 'relu', kernel_initializer='HeUniform')(conv_1)

output_two = Conv2D(filters = 16, kernel_size = (1,1), activation='softmax')(conv_2

flat_ = Flatten() (output_two)

model_two = Model(inputs = vgg_model.input, outputs = flat_)
model_two.summary()
```

Model: "model"

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	[(None, 128, 128, 3)]	0
block1_conv1 (Conv2D)	(None, 128, 128, 64)	1792
block1_conv2 (Conv2D)	(None, 128, 128, 64)	36928
block1_pool (MaxPooling2D)	(None, 64, 64, 64)	0
block2_conv1 (Conv2D)	(None, 64, 64, 128)	73856
block2_conv2 (Conv2D)	(None, 64, 64, 128)	147584
block2_pool (MaxPooling2D)	(None, 32, 32, 128)	0
block3_conv1 (Conv2D)	(None, 32, 32, 256)	295168
block3_conv2 (Conv2D)	(None, 32, 32, 256)	590080
block3_conv3 (Conv2D)	(None, 32, 32, 256)	590080
block3_pool (MaxPooling2D)	(None, 16, 16, 256)	0
block4_conv1 (Conv2D)	(None, 16, 16, 512)	1180160
block4_conv2 (Conv2D)	(None, 16, 16, 512)	2359808
block4_conv3 (Conv2D)	(None, 16, 16, 512)	2359808
block4_pool (MaxPooling2D)	(None, 8, 8, 512)	0
block5_conv1 (Conv2D)	(None, 8, 8, 512)	2359808
block5_conv2 (Conv2D)	(None, 8, 8, 512)	2359808
block5_conv3 (Conv2D)	(None, 8, 8, 512)	2359808
block5_pool (MaxPooling2D)	(None, 4, 4, 512)	0
conv2d (Conv2D)	(None, 1, 1, 2048)	16779264
conv2d_1 (Conv2D)	(None, 1, 1, 2048)	4196352
conv2d_2 (Conv2D)	(None, 1, 1, 16)	32784
flatten (Flatten)	(None, 16)	0
=====		
Total params: 35,723,088		
Trainable params: 21,008,400		
Non-trainable params: 14,714,688		

In [14]: filepath="model_save/best_model_two.hdf5"

```

checkpoint = ModelCheckpoint(filepath=filepath, monitor='val_accuracy', verbose=1,
earlystop = EarlyStopping(monitor='val_accuracy', min_delta=0.01, patience=2, verbo
tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir= ("logs/fits/"+dateti
                    histogram_freq=1,write_graph=True)

call_back_list = [ earlystop, checkpoint, tensorboard_callback]

```

```

In [15]: model_two.compile(optimizer='adam', loss='categorical_crossentropy',metrics=['accu

model_two.fit( train_generator,steps_per_epoch = train_steps_per_epoch,epochs=3,
                validation_data = validation_generator, validation_steps =
                callbacks=[call_back_list])

```

```

Epoch 1/3
563/563 [=====] - 262s 447ms/step - loss: 1.6402 - accura
cy: 0.5293 - val_loss: 1.2483 - val_accuracy: 0.6107

```

```

Epoch 00001: val_accuracy improved from -inf to 0.61067, saving model to model_sav
e\best_model_two.hdf5

```

```

Epoch 2/3
563/563 [=====] - 162s 286ms/step - loss: 1.1337 - accura
cy: 0.6476 - val_loss: 1.1635 - val_accuracy: 0.6364

```

```

Epoch 00002: val_accuracy improved from 0.61067 to 0.63642, saving model to model_
save\best_model_two.hdf5

```

```

Epoch 3/3
563/563 [=====] - 160s 284ms/step - loss: 1.0044 - accura
cy: 0.6868 - val_loss: 1.1059 - val_accuracy: 0.6607

```

```

Epoch 00003: val_accuracy improved from 0.63642 to 0.66075, saving model to model_
save\best_model_two.hdf5

```

```

Out[15]: <keras.callbacks.History at 0x21f003adfa0>

```

```

In [15]: %tensorboard --logdir logs/fits

```


```

Reusing TensorBoard on port 6006 (pid 11692), started 3:18:56 ago. (Use '!kill 116
92' to kill it.)


```


model_two Tensorboard accuracy and loss

epoch_accuracy
tag: epoch_accuracy

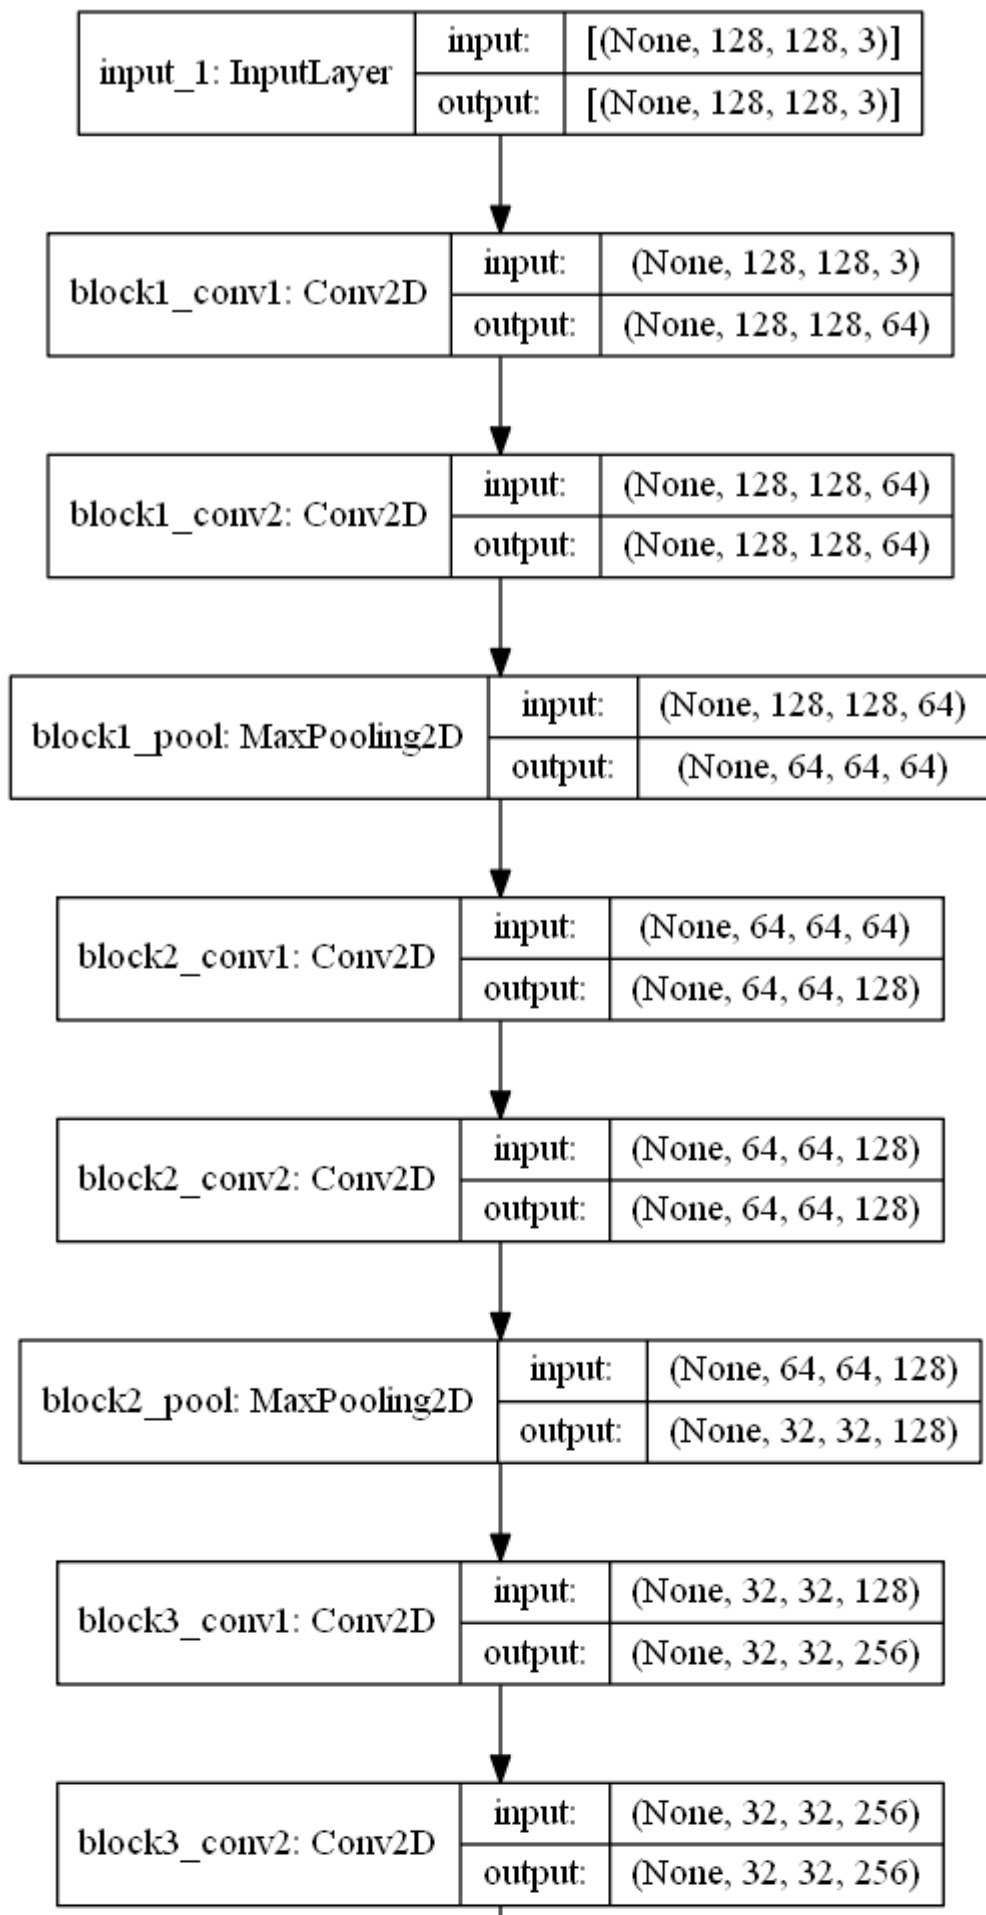


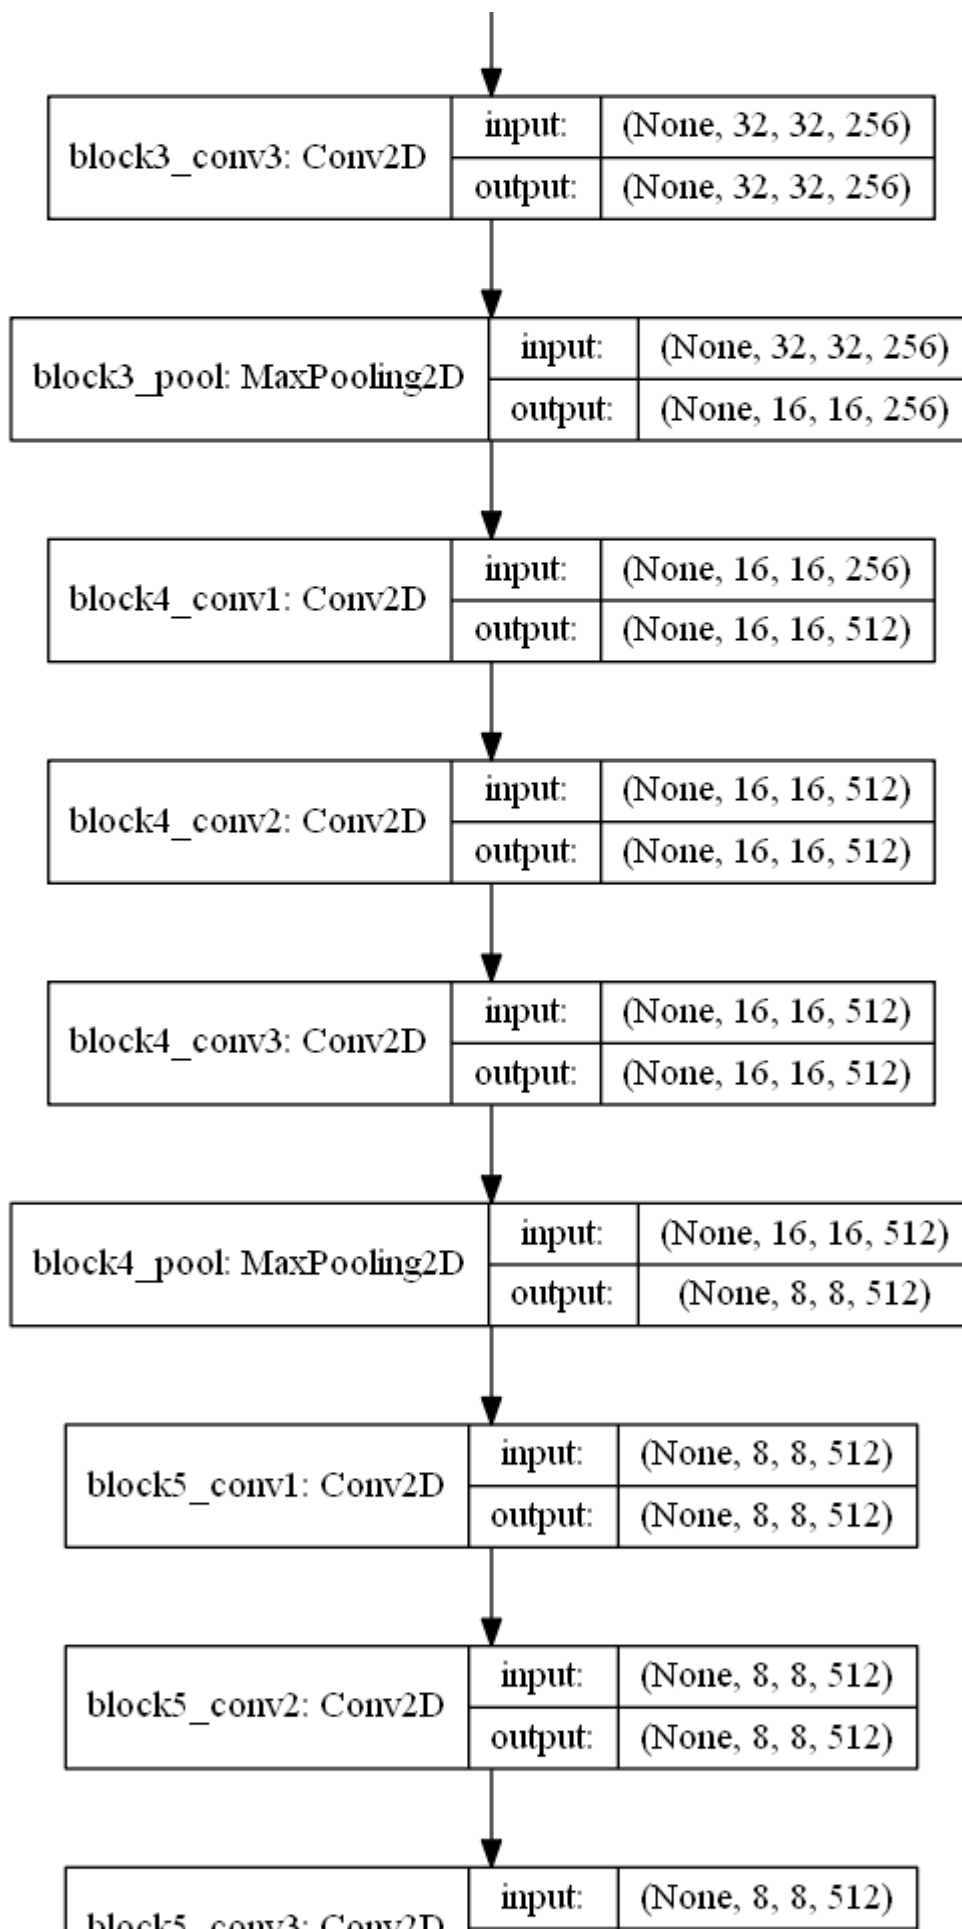
epoch_loss
tag: epoch_loss

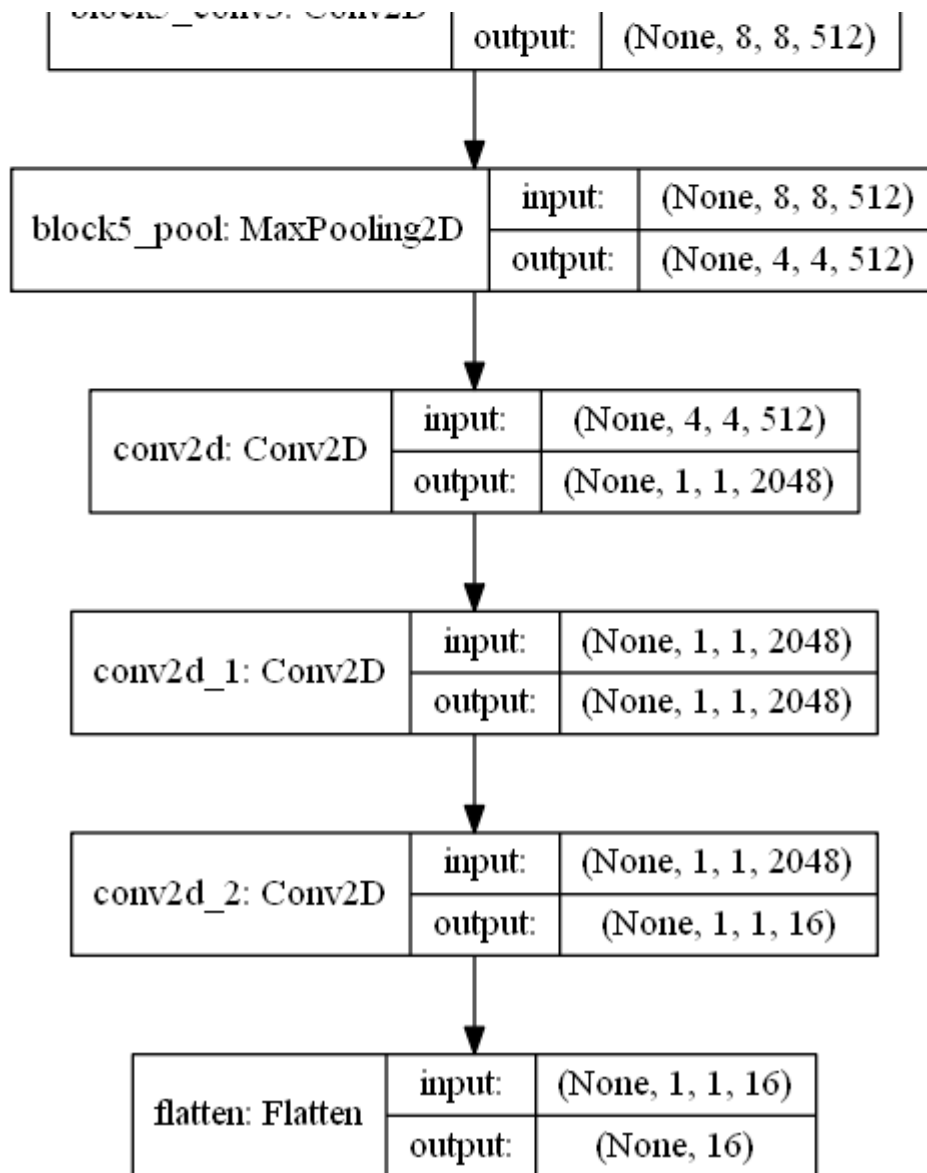


```
In [36]: plot_model(model_two, to_file='model_two.png', show_shapes=True)
```

Out[36]:







Model-3

1. Use same network as Model-2 'INPUT --> VGG-16 without Top layers(FC) --> 2 Conv Layers identical to FC --> Output Layer' and train only Last 6 Layers of VGG-16 network, 2 Conv layers identical to FC layers, 1 output layer.

```
In [30]: tf.keras.backend.clear_session()
```

```
In [31]: #https://blog.keras.io/building-powerful-image-classification-models-using-very-lit
for layer in vgg_model.layers[-6:]:
    layer.trainable = True

for layer in vgg_model.layers[:-6]:
    layer.trainable = False

input_from_vgg = vgg_model.output
```

```
conv_1_ = Conv2D(filters = 2048, kernel_size = (4,4), padding = 'valid',  
                 activation = 'relu', kernel_initializer='HeUniform') (input_from_vgg  
  
conv_2_ = Conv2D(filters = 2048, kernel_size = (1,1), padding = 'valid',  
                 activation = 'relu', kernel_initializer='HeUniform') (conv_1_)  
  
output_three = Conv2D(16, kernel_size = (1,1), activation='softmax')(conv_2_)  
  
flat_3 = Flatten() (output_three)  
  
model_three = Model(inputs = vgg_model.input, outputs = flat_3)  
model_three.summary()
```

Model: "model"

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	[(None, 128, 128, 3)]	0
block1_conv1 (Conv2D)	(None, 128, 128, 64)	1792
block1_conv2 (Conv2D)	(None, 128, 128, 64)	36928
block1_pool (MaxPooling2D)	(None, 64, 64, 64)	0
block2_conv1 (Conv2D)	(None, 64, 64, 128)	73856
block2_conv2 (Conv2D)	(None, 64, 64, 128)	147584
block2_pool (MaxPooling2D)	(None, 32, 32, 128)	0
block3_conv1 (Conv2D)	(None, 32, 32, 256)	295168
block3_conv2 (Conv2D)	(None, 32, 32, 256)	590080
block3_conv3 (Conv2D)	(None, 32, 32, 256)	590080
block3_pool (MaxPooling2D)	(None, 16, 16, 256)	0
block4_conv1 (Conv2D)	(None, 16, 16, 512)	1180160
block4_conv2 (Conv2D)	(None, 16, 16, 512)	2359808
block4_conv3 (Conv2D)	(None, 16, 16, 512)	2359808
block4_pool (MaxPooling2D)	(None, 8, 8, 512)	0
block5_conv1 (Conv2D)	(None, 8, 8, 512)	2359808
block5_conv2 (Conv2D)	(None, 8, 8, 512)	2359808
block5_conv3 (Conv2D)	(None, 8, 8, 512)	2359808
block5_pool (MaxPooling2D)	(None, 4, 4, 512)	0
conv2d (Conv2D)	(None, 1, 1, 2048)	16779264
conv2d_1 (Conv2D)	(None, 1, 1, 2048)	4196352
conv2d_2 (Conv2D)	(None, 1, 1, 16)	32784
flatten (Flatten)	(None, 16)	0
=====		
Total params: 35,723,088		
Trainable params: 30,447,632		
Non-trainable params: 5,275,456		

In [32]: filepath="model_save/best_model_three.hdf5"

```

checkpoint = ModelCheckpoint(filepath=filepath, monitor='val_accuracy', verbose=1,
earlystop = EarlyStopping(monitor='val_accuracy', min_delta=0.001, patience=2, verb
tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir= ("logs/fits/"+dateti
                    histogram_freq=1,write_graph=True)

call_back_list = [ earlystop, checkpoint, tensorboard_callback]

```

```

In [33]: model_three.compile(optimizer='adam', loss='categorical_crossentropy',metrics=['acc

model_three.fit( train_generator,steps_per_epoch = train_steps_per_epoch,epochs=5,
                validation_data = validation_generator, validation_steps =
                callbacks=[call_back_list])

```

```

Epoch 1/5
563/563 [=====] - 273s 483ms/step - loss: 2.7728 - accuracy: 0.0619 - val_loss: 2.7727 - val_accuracy: 0.0619

```

```

Epoch 00001: val_accuracy improved from -inf to 0.06192, saving model to model_save\best_model_three.hdf5

```

```

Epoch 2/5
563/563 [=====] - 279s 496ms/step - loss: 2.7727 - accuracy: 0.0607 - val_loss: 2.7730 - val_accuracy: 0.0593

```

```

Epoch 00002: val_accuracy did not improve from 0.06192

```

```

Epoch 3/5
563/563 [=====] - 278s 493ms/step - loss: 2.7727 - accuracy: 0.0627 - val_loss: 2.7730 - val_accuracy: 0.0593

```

```

Epoch 00003: val_accuracy did not improve from 0.06192

```

```

Epoch 00003: early stopping

```

```

Out[33]: <keras.callbacks.History at 0x2204d567eb0>

```

```

In [12]: %tensorboard --logdir logs/fits

```


```

Reusing TensorBoard on port 6006 (pid 11692), started 15:28:36 ago. (Use '!kill 11692' to kill it.)


```


model_three Tensorboard accuracy and loss

epoch_accuracy
tag: epoch_accuracy

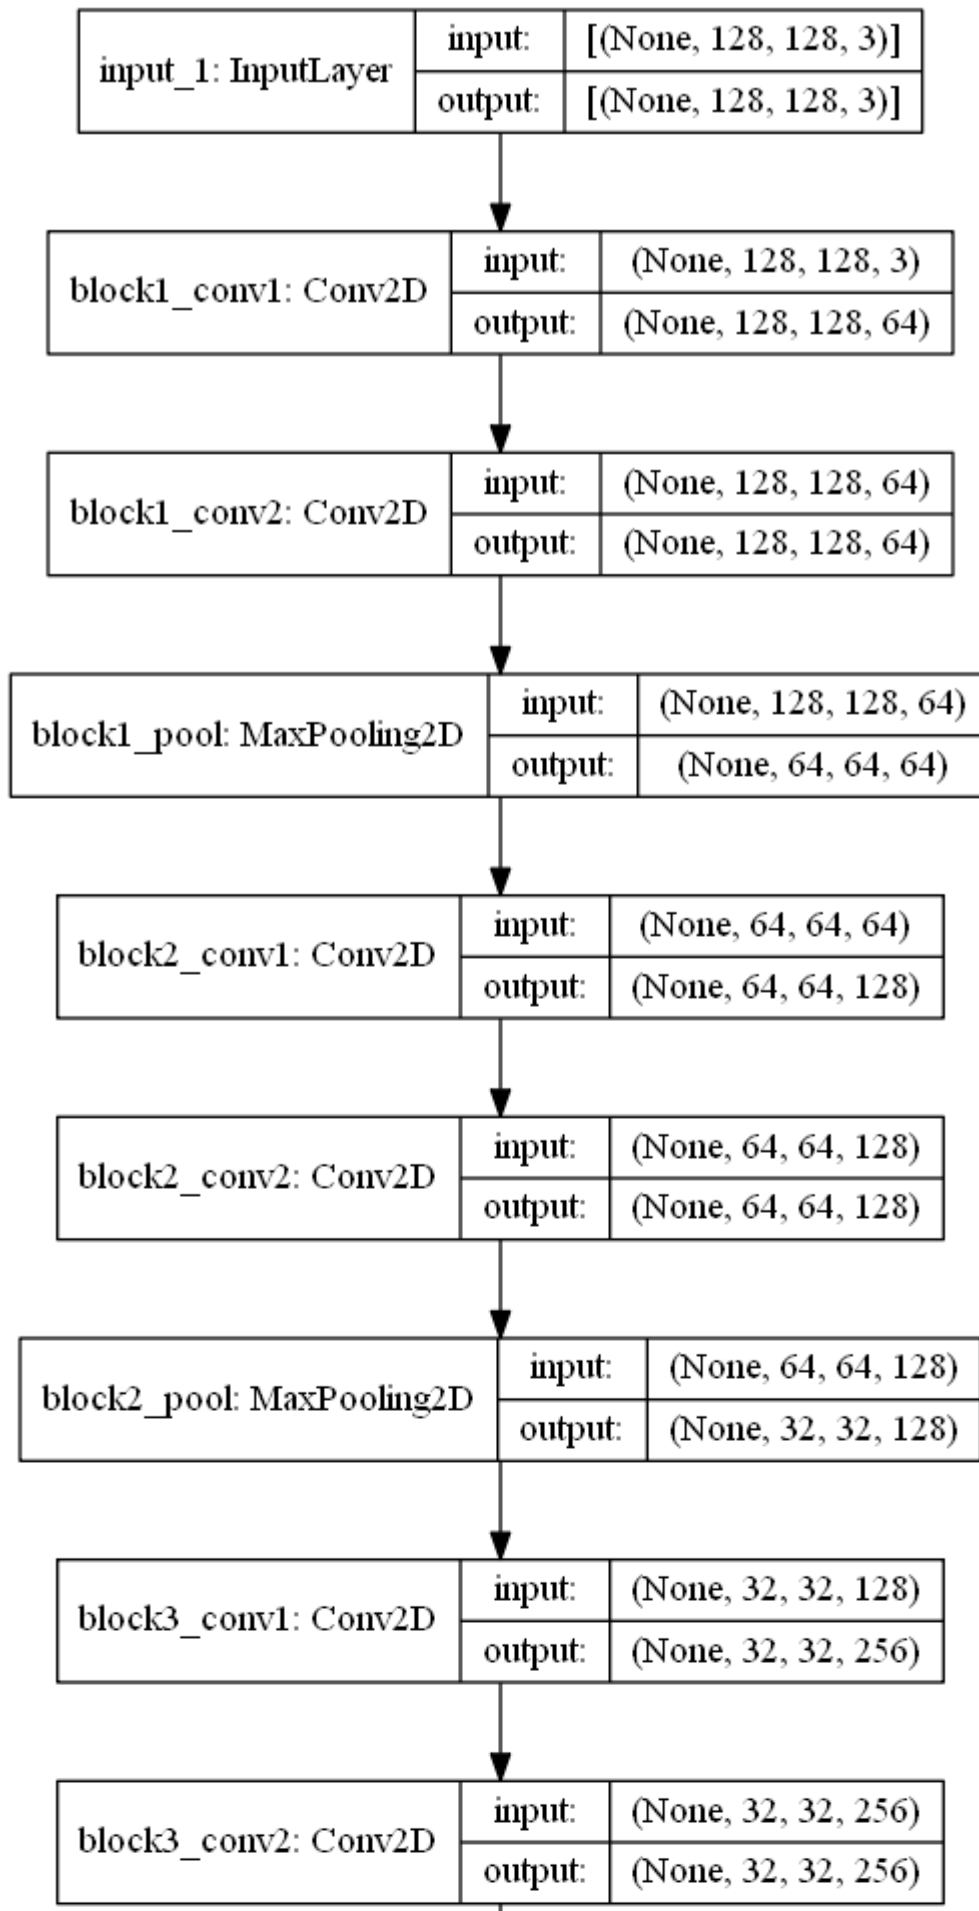


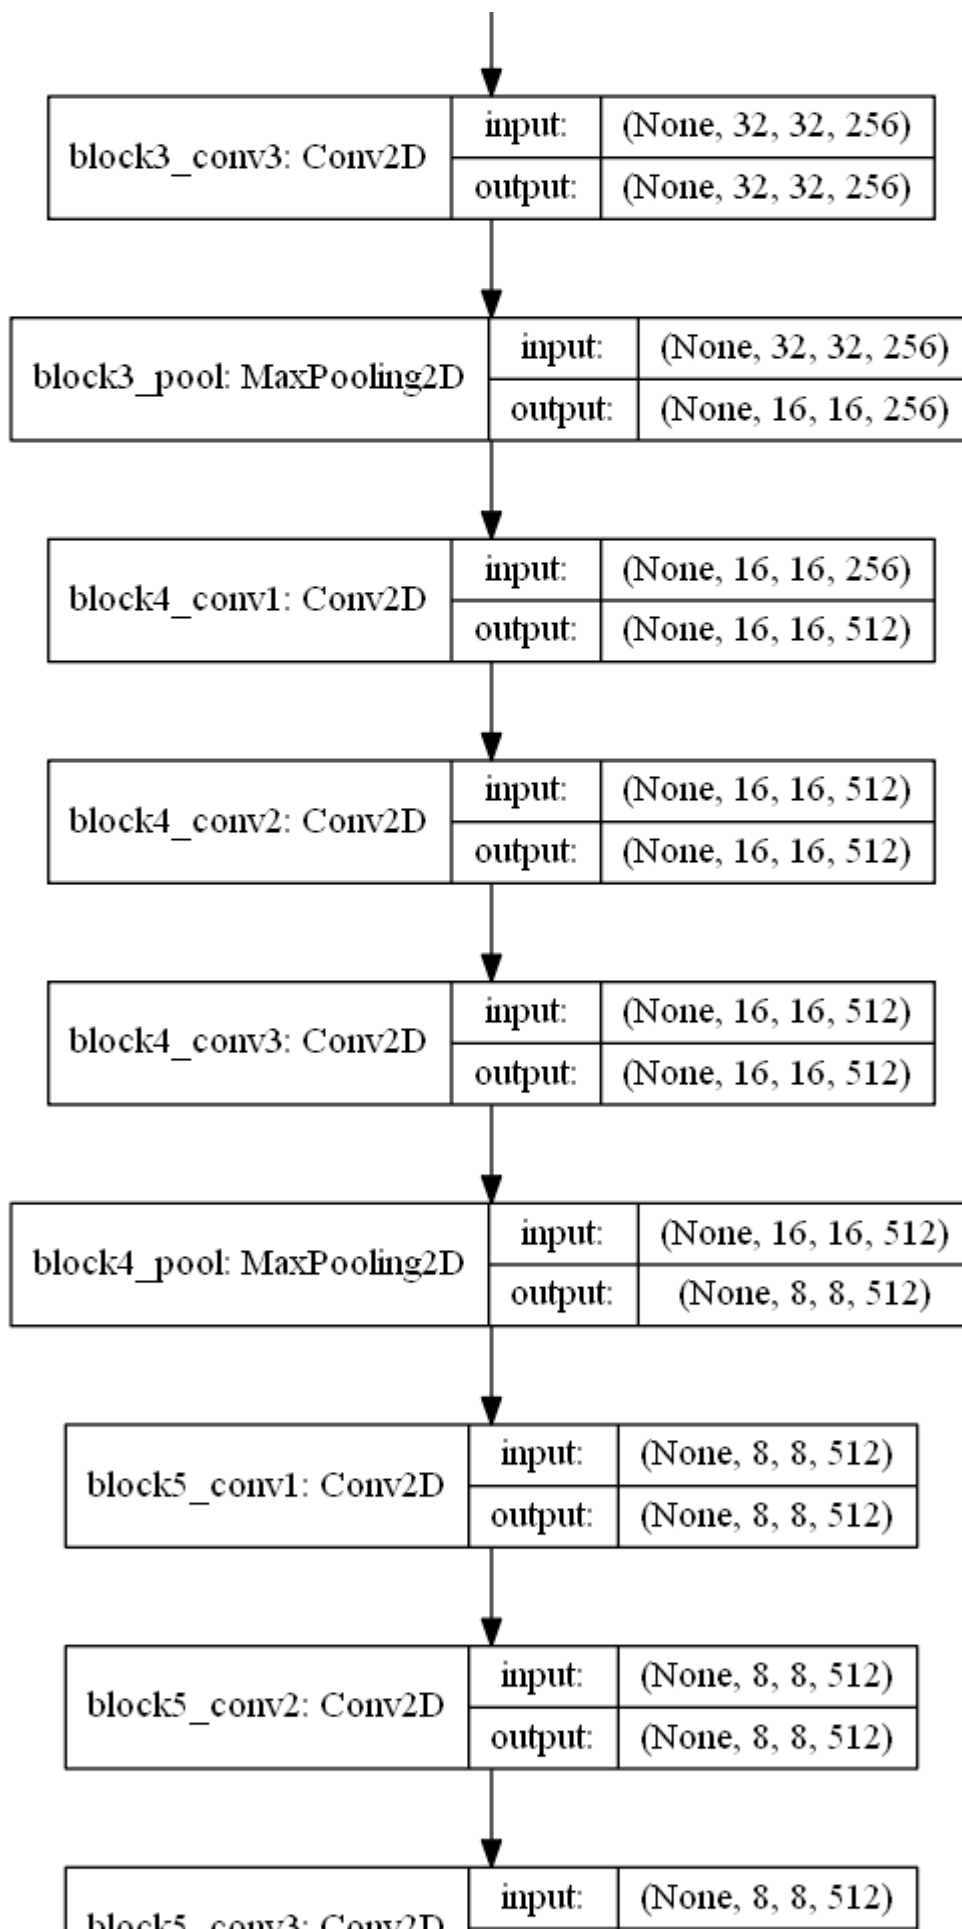
epoch_loss
tag: epoch_loss

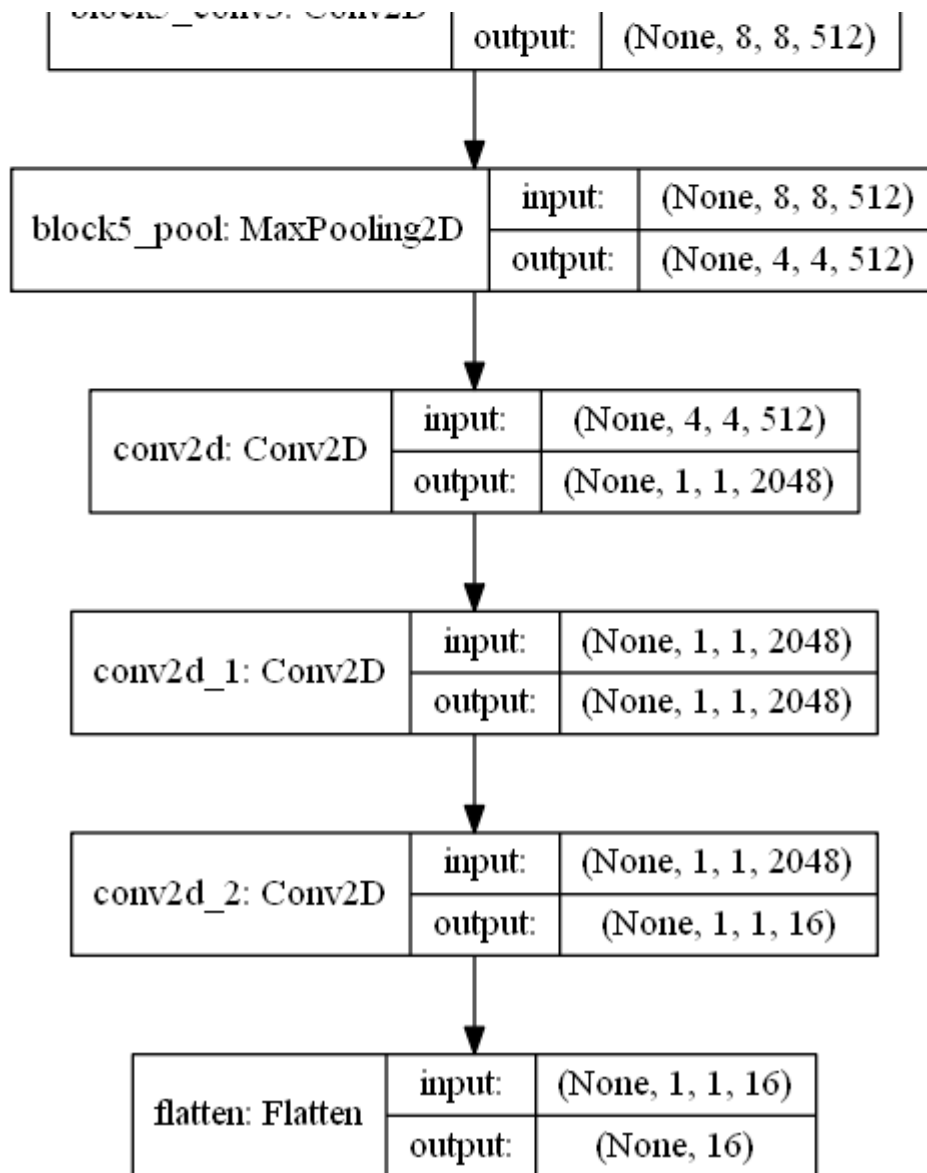


```
In [35]: plot_model(model_three, to_file='model_three.png', show_shapes=True)
```

Out[35]:







Please write your observations or a brief summary of the results that you get after performing transfer learning with reference to model1, model2 and model3

```

In [34]: # http://zetcode.com/python/prettytable/

from prettytable import PrettyTable

x = PrettyTable()
x.field_names = ["Model", "Epochs", "val_accuracy"]

x.add_row(["model_one", "3", "0.6526"])
x.add_row(["model_two", "3", "0.6607"])
x.add_row(["model_three", "3", "0.0619"])

print(x)

```

Model	Epochs	val_accuracy
model_one	3	0.6526
model_two	3	0.6607
model_three	3	0.0619

Observations:

1. model_one only had fully connected layers after vgg-16 model as trainable layers and also had the least trainable parameters among all three models (2,919,312).
2. model_two used convolution layers instead of fully connected layers after vgg-16 model as trainable layers and had the second most trainable parameters among all three models (21,008,400).
3. model_three used the same architecture of model_two, but also had the last 6 layers in vgg-16 model to be set as trainable. Because of training these additional layers, it had the highest number of trainable parameters among all three models (30,447,632).
4. model_one took the least time to train because it had the least number of trainable parameters while model_two took longer and model_three took the longest time to train as they had the higher number of trainable parameters.
5. Since model_three has 9 layers and high number of parameters to train, it needs significant number of epochs to get higher accuracy. This is the reason why model_three only had 6% accuracy after three epochs.
6. Since model_one and model_two only had 3 layers to be trained and used the pre-trained weights from vgg-16 model, they have more than 60% accuracy after 3 epochs.